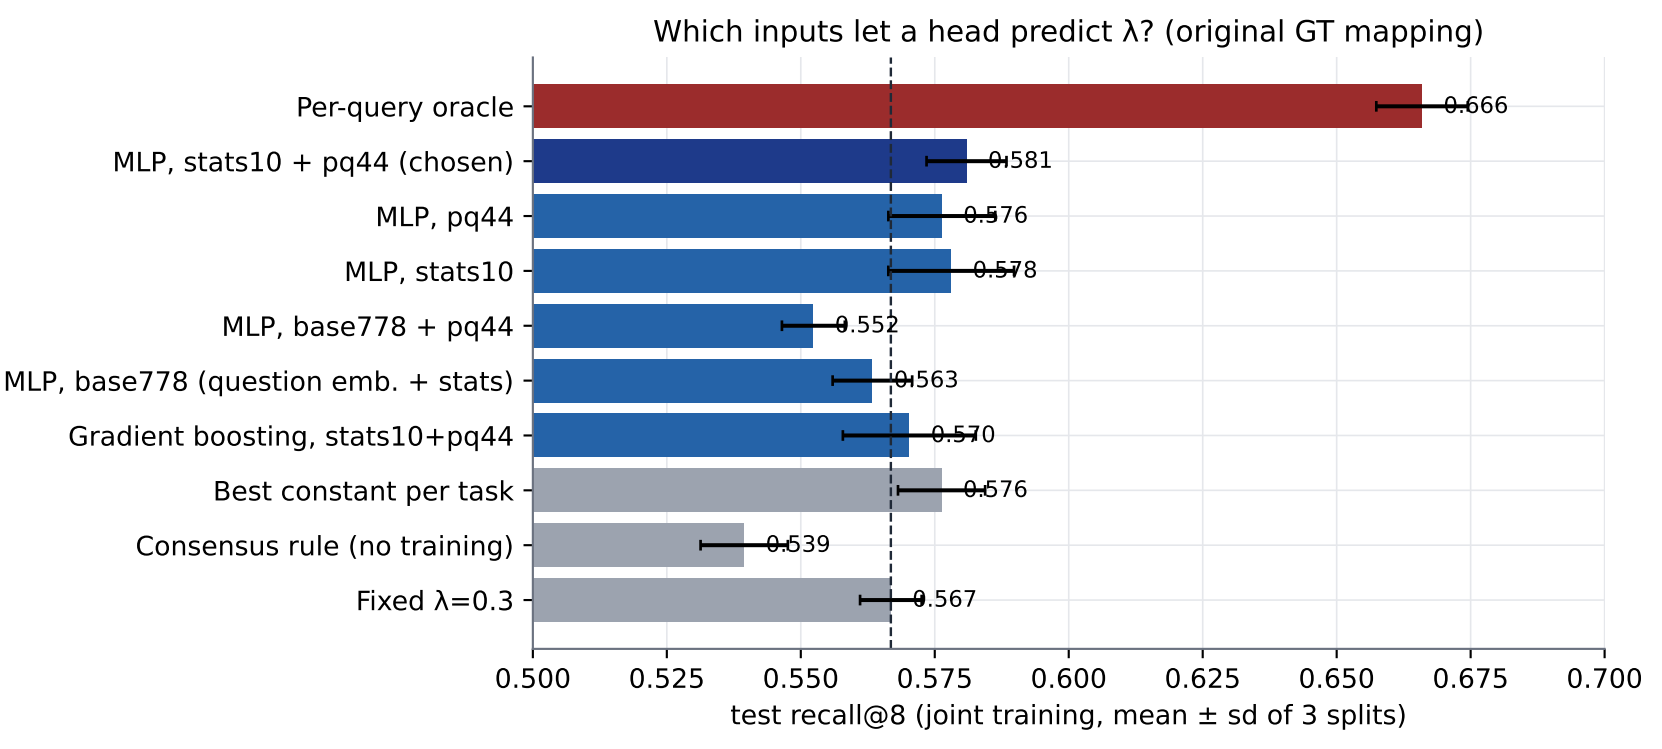
Which inputs let a head predict λ? (original GT mapping)
Per-query oracle
MLP, stats10 + pq44 (chosen)
MLP, pq44
MLP, stats10
MLP, base778 + pq44
MLP, base778 (question emb. + stats)
Gradient boosting, stats10+pq44
Best constant per task
Consensus rule (no training)
Fixed λ=0.3
0.666
0.581
0.576
0.578
0.552
0.563
0.570
0.576
0.539
0.567
0.500
0.525
0.550
0.575
0.600
0.625
0.650
0.675
0.700
test recall@8 (joint training, mean ± sd of 3 splits)